

```
In [1]:
```

```
import os
print(os.getcwd())
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import GridSearchCV
%matplotlib inline
```

```
C:\Users\arun_
```

```
In [2]:
```

```
df = pd.read_csv('car_evaluation.csv', header = None)
```

```
In [3]:
```

```
df.head()
```

```
Out[3]:
```

0	1	2	3	4	5	6	
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [4]:
```

```
col_names = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
df.columns = col_names
col_names
```

```
Out[4]:
```

```
['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
```

```
In [5]:
```

```
df.head()
```

```
Out[5]:
```

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [6]:
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   buying      1728 non-null   object 
 1   maint       1728 non-null   object 
 2   doors       1728 non-null   int64  
 3   persons     1728 non-null   int64  
 4   lug_boot    1728 non-null   object 
 5   safety      1728 non-null   object 
 6   class       1728 non-null   object 
```

```
buying    1728 non-null object
1  maint    1728 non-null object
2  doors    1728 non-null object
3  persons  1728 non-null object
4  lug_boot 1728 non-null object
5  safety   1728 non-null object
6  class    1728 non-null object
dtypes: object(7)
memory usage: 94.6+ KB
```

In [7]:

```
for i in col_names:
    print(df[i].value_counts())
```

```
vhigh    432
high     432
med      432
low      432
Name: buying, dtype: int64
vhigh    432
high     432
med      432
low      432
Name: maint, dtype: int64
2        432
3        432
4        432
5more   432
Name: doors, dtype: int64
2        576
4        576
more    576
Name: persons, dtype: int64
small   576
med     576
big     576
Name: lug_boot, dtype: int64
low     576
med     576
high    576
Name: safety, dtype: int64
unacc  1210
acc    384
good   69
vgood  65
Name: class, dtype: int64
```

In [8]:

```
df.shape
```

Out[8]:

```
(1728, 7)
```

In [9]:

```
X = df.drop(['class'], axis = 1)
y = df['class']
```

In [10]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=42)
```

In [11]:

```
from sklearn.preprocessing import OrdinalEncoder
enc = OrdinalEncoder()
X_train = enc.fit_transform(X_train)
X_test = enc.transform((X_test))
```

## Gini index as criterion

In [12]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [13]:

```
clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=42)
clf_gini.fit(X_train, y_train)
```

Out[13]:

```
DecisionTreeClassifier(max_depth=3, random_state=42)
```

In [14]:

```
y_pred = clf_gini.predict(X_test)
```

In [15]:

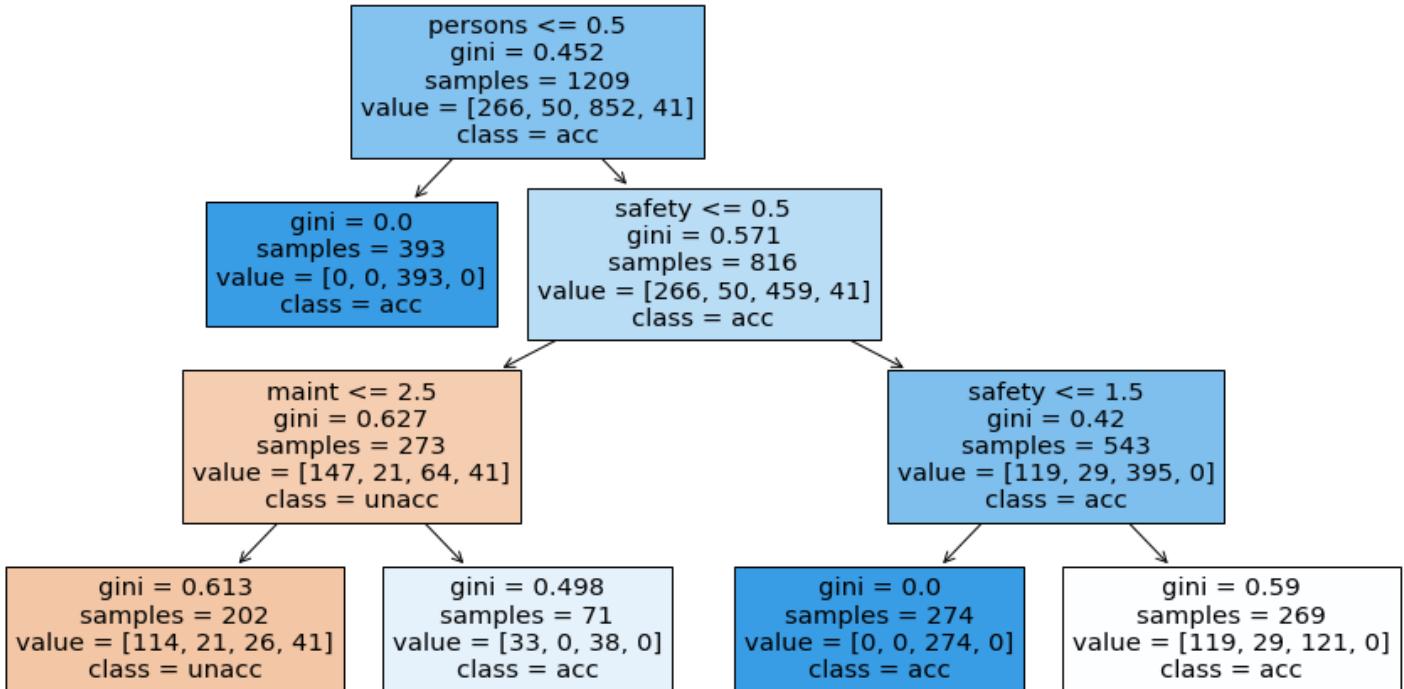
```
from sklearn.metrics import accuracy_score

print(f'Model with gini index gives an accuracy of: {accuracy_score(y_true=y_test, y_pred=y_pred)}')
```

Model with gini index gives an accuracy of: 0.7572254335260116

In [16]:

```
from sklearn import tree
plt.figure(figsize=(15,8))
tree.plot_tree(clf_gini,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()
```



In [17]:

```
# Check for underfitting

print(f'Training set score: {clf_gini.score(X_train,y_train)}')
```

```
print(f'Test set score: {clf_gini.score(X_test,y_test)}')
```

Training set score: 0.7775020678246485  
Test set score: 0.7572254335260116

## Entropy as criterion

In [18]:

```
clf_entropy = DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=42)
clf_entropy.fit(X_train, y_train)
```

Out[18]:

```
DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=42)
```

In [19]:

```
y_pred = clf_entropy.predict(X_test)
```

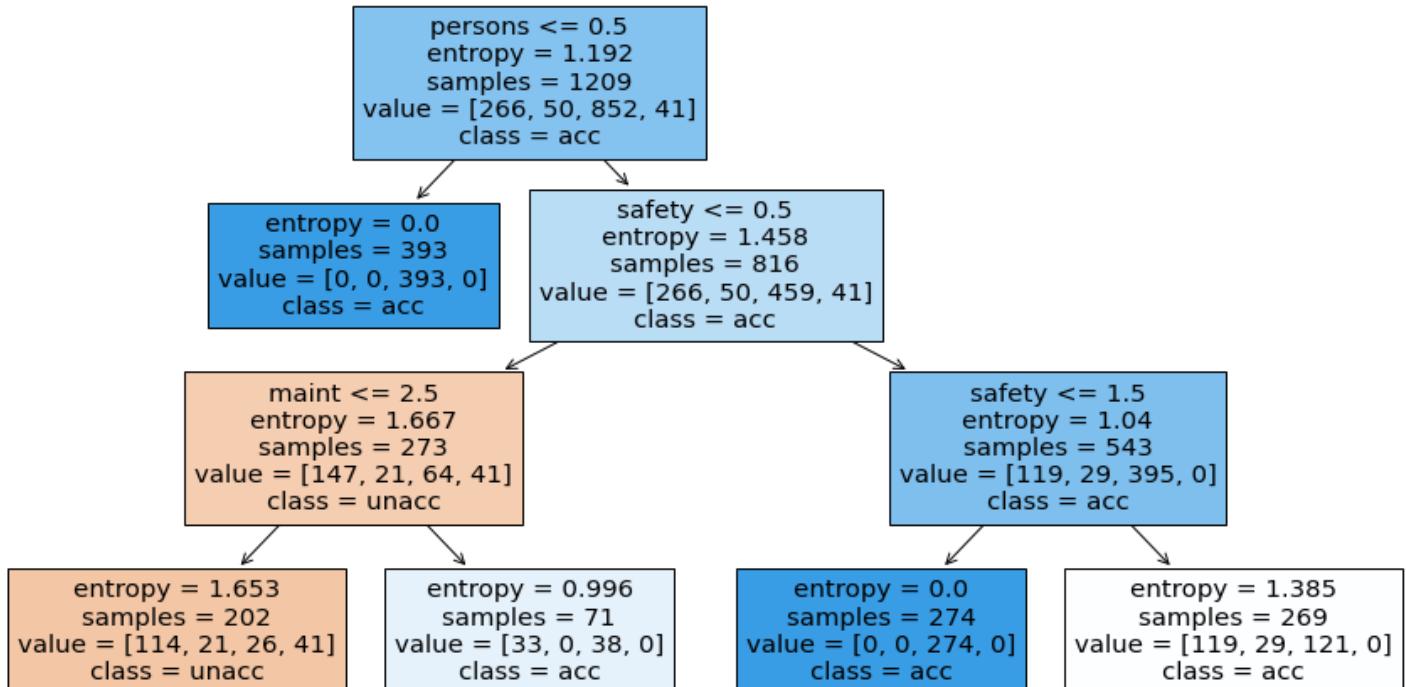
In [20]:

```
from sklearn.metrics import accuracy_score
print(f'Model with gini index gives an accuracy of: {accuracy_score(y_test, y_pred)}')
```

Model with gini index gives an accuracy of: 0.7572254335260116

In [21]:

```
plt.figure(figsize=(15,8))
tree.plot_tree(clf_entropy,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()
```



In [22]:

```
# Check for underfitting
print(f'Training set score: {clf_entropy.score(X_train,y_train)}')
print(f'Test set score: {clf_entropy.score(X_test,y_test)}')
```

```
Training set score: 0.7572254335260116
Test set score: 0.7572254335260116
```

In [23]:

```
from sklearn.metrics import confusion_matrix, classification_report
cm = confusion_matrix(y_test, y_pred)
```

In [24]:

```
print(cm)
```

```
[[ 44   0   74   0]
 [  9   0   10   0]
 [  9   0  349   0]
 [ 24   0   0   0]]
```

In [25]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
acc	0.51	0.37	0.43	118
good	0.00	0.00	0.00	19
unacc	0.81	0.97	0.88	358
vgood	0.00	0.00	0.00	24
accuracy			0.76	519
macro avg	0.33	0.34	0.33	519
weighted avg	0.67	0.76	0.71	519

```
C:\Users\PRATYUSH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\PRATYUSH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\PRATYUSH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

## Cross Validation

In [26]:

```
params_grid = {
    'criterion':['gini','entropy'],
    'max_depth':[3,4,5,6,7,8,9,10]
}
```

In [27]:

```
decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_train,y_train)
```

Out[27]:

```
DecisionTreeClassifier()
```

In [28]:

```
dt_validated = GridSearchCV(estimator=decision_tree, param_grid=params_grid,scoring='accuracy',cv=20)
```

In [29]:

```
%%time  
dt_validated.fit(X_train, y_train)
```

Wall time: 840 ms

Out[29]:

```
GridSearchCV(cv=20, estimator=DecisionTreeClassifier(),  
            param_grid={'criterion': ['gini', 'entropy'],  
                        'max_depth': [3, 4, 5, 6, 7, 8, 9, 10]},  
            scoring='accuracy')
```

In [30]:

```
print(f'Best parameters for decision tree classifier after CV -> {dt_validated.best_params_}')  
print(f'Best score on decision tree classifier after CV -> {dt_validated.best_score_}')
```

Best parameters for decision tree classifier after CV -> {'criterion': 'entropy', 'max\_depth': 10}

Best score on decision tree classifier after CV -> 0.979330601092896

In [31]:

```
print(f'Score on train set of DT classifier before CV -> {decision_tree.score(X_train, y_train)}')  
print(f'Score on test set of DT classifier before CV -> {decision_tree.score(X_test, y_test)}')  
print(f'Score on train set of DT classifier after CV -> {dt_validated.score(X_train, y_train)}')  
print(f'Score on test set of DT classifier after CV -> {dt_validated.score(X_test, y_test)}')
```

Score on train set of DT classifier before CV -> 1.0

Score on test set of DT classifier before CV -> 0.9653179190751445

Score on train set of DT classifier after CV -> 0.9925558312655087

Score on test set of DT classifier after CV -> 0.9595375722543352

In [32]:

```
print('Classification report on train set')  
print(classification_report(y_true=y_train, y_pred=dt_validated.predict(X_train)))
```

Classification report on train set

	precision	recall	f1-score	support
acc	0.98	0.99	0.99	266
good	0.98	0.98	0.98	50
unacc	1.00	0.99	1.00	852
vgood	0.98	1.00	0.99	41
accuracy			0.99	1209
macro avg	0.98	0.99	0.99	1209
weighted avg	0.99	0.99	0.99	1209

In [33]:

```
print('Classification report on test set')  
print(classification_report(y_true=y_test, y_pred=dt_validated.predict(X_test)))
```

Classification report on test set

	precision	recall	f1-score	support
acc	0.92	0.92	0.92	118
good	0.71	0.89	0.79	19
unacc	0.99	0.98	0.99	358
vgood	0.88	0.88	0.88	24
accuracy			0.96	519
macro avg	0.88	0.92	0.89	519
weighted avg	0.96	0.96	0.96	519

In [34]:

```
print('Confusion matrix on train set')
print(confusion_matrix(y_true=y_train, y_pred=dt_validated.predict(X_train)))
```

Confusion matrix on train set

```
[[263  1  2  0]
 [ 0  49  0  1]
 [ 5  0 847  0]
 [ 0  0  0  41]]
```

In [35]:

```
print('Confusion matrix on test set')
print(confusion_matrix(y_true=y_test, y_pred=dt_validated.predict(X_test)))
```

Confusion matrix on test set

```
[[108  7  2  1]
 [ 0 17  0  2]
 [ 6  0 352  0]
 [ 3  0  0  21]]
```

In [39]:

```
best_dt = DecisionTreeClassifier(criterion='entropy', max_depth=9)
```

In [40]:

```
best_dt.fit(X_train,y_train)
```

Out[40]:

```
DecisionTreeClassifier(criterion='entropy', max_depth=9)
```

In [41]:

```
plt.figure(figsize=(15,8))
tree.plot_tree(best_dt,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()
```

